

NWEN302

Lab 3

Olivia Fletcher
300534281
fletcholiv

REPORT

Introduction

In this investigation we will be utilizing a tool known as Software-Defined Networking (SDN) where its prime functionality is using software-based controllers or API's to ensure a more efficient network performance. The use of SDN resembles cloud commuting as it enables dynamic configuration and monitoring duties unlike a traditional way of network management.



Design

We will be utilizing Python and implementing network simulation through CORE gui. In this investigation we will be learning the concepts behind SDN networks (Software-Defined Networking) and its available operations such as Ryu and OpenFlow. The Ryu will be acting as a controller to the CORE network simulation. Ryu will be extending the use of OpenFlow as a means to implement the Python code to orchestrate the network simulation.

We will be adding to the provided `simple_switch13.py` code to complete task1 ;

- Block traffic between host 2 and host 3

Due to the source and destination mac addresses being provided I just need to do a simple if statement that checks if the source and/or destination is host2 and/or host 3 and blocks the communications.

```

nwen302_lab3.py X
C: > Users > OEM > Documents > Uni > NWEN302 > Lab3 > nwen302_lab3.py
1  # Simple Switch 13 Psuedo Code - Task 1
2  # The goal is to block traffic between host2 and host3 via python code on the
3  # running network simulation.
4
5  # The provided python code includes the source and destination MAC address -
6  # and the ability to route traffic between hosts. Before the communication -
7  # takes places in the python code I need to include a couple if statements -
8  # to check the host details and drop the communications.
9  if [source == "host 2 mac address" and destination == "host 3 mac address"]
10     if [source "host 3 mac address" and destination == "host 2 mac address"]
11         # Print to terminal
12         print "Blocked communications"
13         # empty return statement to exit loop
14         return

```

To complete task 2;

- Count all traffic communicating with host 1

The first block follows similar to task 1 where the second block of code is after making changes when I encountered some errors/issues (I go into detail later on)

```

16 # Simple Switch 13 Psuedo Code - Task 2
17 # The goal is to count all traffic communicating with host1
18
19 # The provided python code includes the source and destination MAC address -
20 # and the ability to route traffic between hosts. Before the communication -
21 # takes places in the python code I need to include an if statement that -
22 # checks if the source or destination address is from host 1 and to count -
23 # the traffic
24
25 # Global declaration ;
26 int num = 0
27
28 # Fetch the global field num
29 Global num
30 if [source == "host 1 mac address" OR destination == "host 1 mac address"]
31     # Add to count variable
32     num += 1
33     # Print to terminal
34     print "Blocked communications"
35     # empty return statement to exit loop
36     return
37
38
39 # OR, after fixing code
40 within the def __init__() function
41     add a int counter, int initnum and boolean init variable for later use
42
43 # Implement the following functions from the traffic_monitor.py
44 def _state_change_handler()
45 def _monitor()
46 def _request_stats()
47 def _port_stats_reply_handler()
48 using the provided sorted(body, key=attrgetter('port_no')) set to variable stats for use
49
50 if the self.init (true)
51     set the first element of the stats dictionary of the rx_packets and tx_packets to the initnum variable
52     print the number to the terminal screen for testing
53     set the init to false
54
55 set the count variable to equal the first elements (host 1) to equal the traffic - the initnum variable
56 (we have to do this to count because the way the ryu is set up it holds the traffic history so the count is inaccurate)
57 print the count to the terminal
58

```

To complete task 3;

- Count traffic communications from each host and once the number has hit a MAX_COUNT to block communications from that host for 60 seconds.

This task we will be combining task 1 and task 2 by implementing a blocking algorithm after a given host's communications count has hit a maximum amount. The below 2 images are for my first iteration of my attempt for task 3.

```
38 # Simple Switch 13 Psuedo Code - Task 3
39 # The goal is to count all traffic from each host and to block communication -
40 # after a set MAX_COUNT number and keep it blocked until 60 seconds
41
42 # The provided python code includes the source and destination MAC address -
43 # and the ability to route traffic between hosts. Before the communication -
44 # takes places in the python code I need to include an if statement that -
45 # checks if the source or destination address is from host 1 and to count -
46 # the traffic
47
48 # Global declarations ;
49 int MAX_COUNT = 10
50
51 # Maps/dictionaries for counting and blocking
52 Counter Dictionary = {"host 1 mac address" : 0,
53                       "host 2 mac address" : 0,
54                       "host 3 mac address" : 0}
55
56 Blocking Dictionary = {"host 1 mac address" : False,
57                        "host 2 mac address" : False,
58                        "host 3 mac address" : False}
59
60 # Before entering the main packet sending/receiving function create a method for thread closing
61 def timerChecking(mac address)
62     # Set the current mac address within the counter dictionary to false to reset blocking
63     Counter[mac address] == false
64
65 # Now go inside the main packet sending method '_packet_in_handler'
66 # Fetch the global field num
67 Global num
68 Global Counter
69 Global Blocking
```

```
70
71 if the source from where the traffic is originating from is in the self.mac_addresses
72     add the element at the source addresses position in the Counter dictionary to add 1
73
74 loop through the key and value pairs in the Blocking dictionary
75     Check if the key is equal to the source and the value is assigned as True
76     # Block the communication
77     return
78
79 Finally, check that the Count dictionary element at the source addresses position is equal
80 to the MAX_COUNT, int 10 variable
81     if so, set the Blocking dictionarys element at the source address position to true
82
83     execute the thread which will continue until 60 seconds and once complete, call the timerChecking
84     method to unblock
85
```

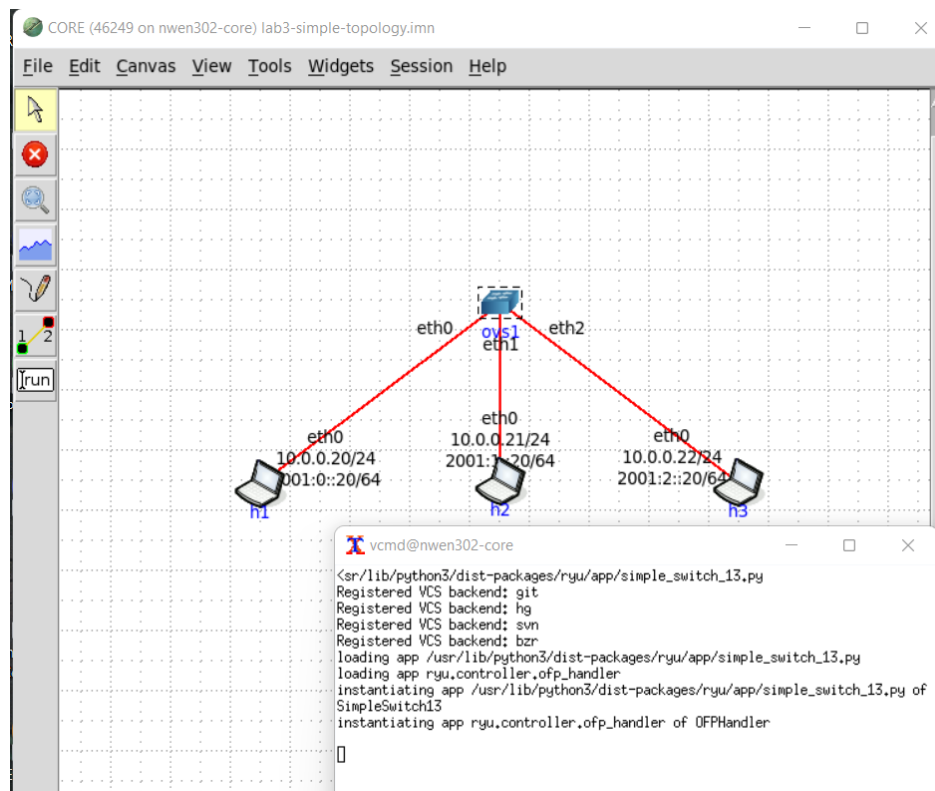
- My second attempt at task 3 after making some changes; (I discuss why later on)

```
81
82 # OR, alternative
83 within the def __init__() function
84     add the following for later use:
85     int initnumhost1,
86     int initnumhost2,
87     int initnumhost3,
88     int initnum
89     bool init
90
91     addrCount dictionary which contains each hosts mac address with an assigned int value
92     addrTimer dictionary which contains each hosts mac address with an assigned boolean value
93
94 # Implement the following functions from the traffic_monitor.py
95 def _state_change_handler()
96 def _monitor()
97 def _request_stats()
98 def _port_stats_reply_handler()
99 using the provided sorted(body, key=attrgetter('port_no')) set to variable stats for use
100
101 if self.initnumhost1 (true)
102     set the first[0] element of the stats dictionary of the rx_packets and tx_packets to the initnum variable
103     print the number to the terminal screen for testing
104     set the init to false
105
106 if self.initnumhost2 (true)
107     set the second[1] element of the stats dictionary of the rx_packets and tx_packets to the initnum variable
108     print the number to the terminal screen for testing
109     set the init to false
110
111 if self.initnumhost3 (true)
112     set the third[2] element of the stats dictionary of the rx_packets and tx_packets to the initnum variable
113     print the number to the terminal screen for testing
114     set the init to false
115
116 set the count variable to equal the first elements (host 1) to equal the traffic - the hosts initnum variable
117 (we have to do this to count because the way the ryu is set up it holds the traffic history so the count is inaccurate)
118 print the count to the terminal
119
120 do the above 3 lines again but for hosts 2 and 3
```

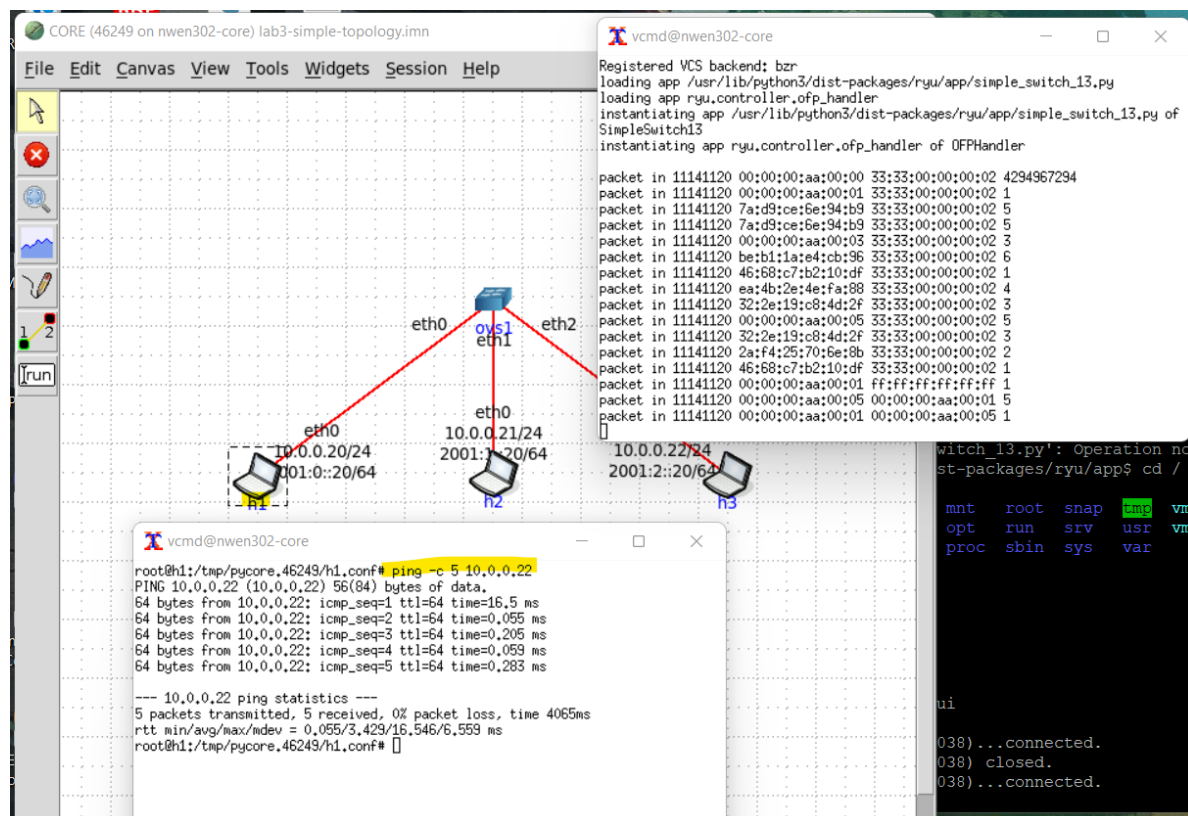
Procedures

I created a simple network topology which consists of 3 hosts connected to a switch router. The switch is not like a regular layer 2 device but a layer 3 Open vSwitch which will allow for more capabilities such as the use of OpenFlow controller.

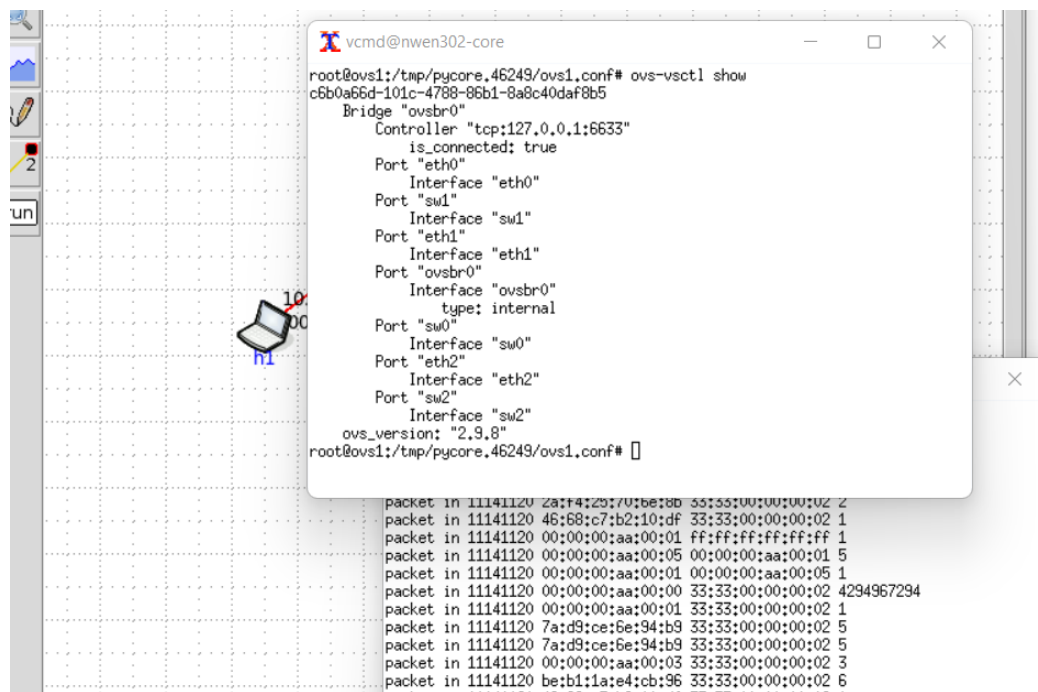
Before making any adjustments to the provided python code I have run it through the switches terminal and executed a few commands to test preemptively. Below is a screenshot showing the code running through the switches terminal.



When ping from host 1 to host 3 the switches terminal appears as below; showing that it foresees every communication/traffic between hosts.



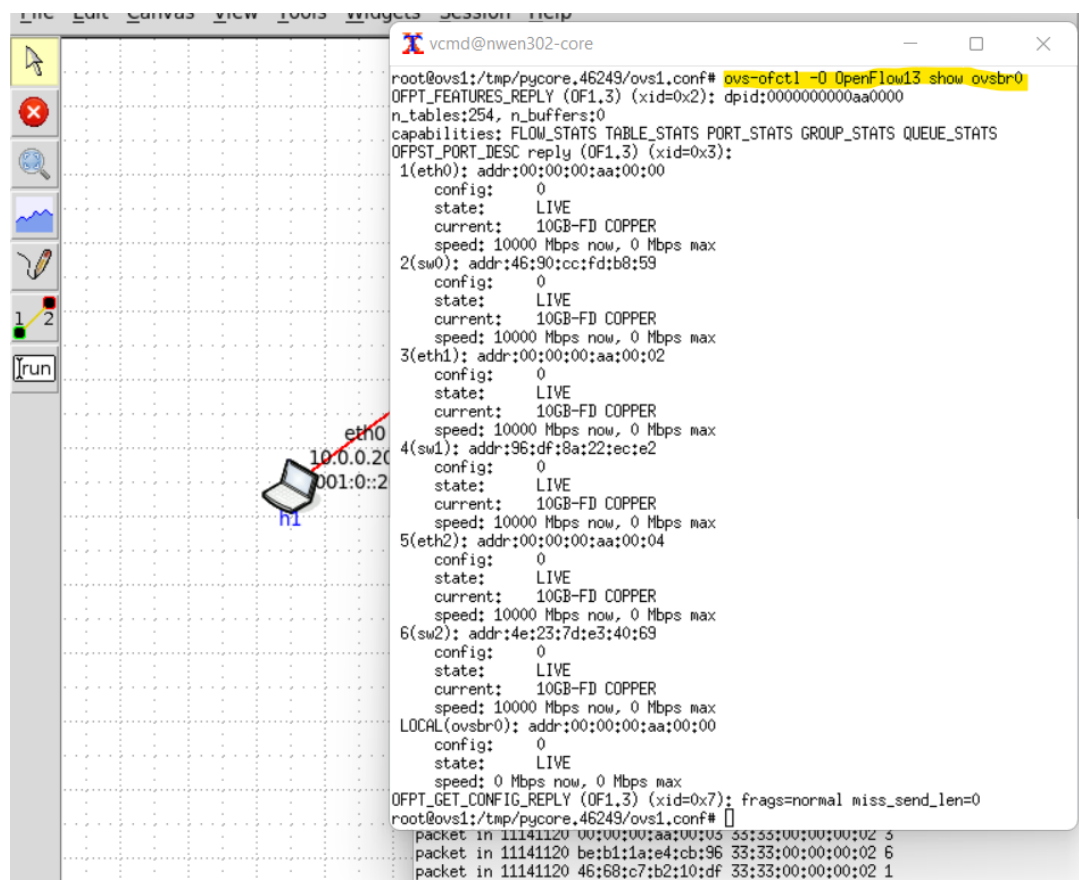
The command 'ovs-vsctl show' prints an overview of the switch database configuration while 'ovs-ofctl -O OpenFlow13 show ovsbr0' shows an overview of the switches OpenFlow configuration.



```

vcmd@nwen302-core
root@ovs1:/tmp/pycore.46249/ovs1.conf# ovs-vsctl show
c6b0a66d-101c-4788-86b1-8a8c40daf8b5
Bridge "ovsbr0"
  Controller "tcp:127.0.0.1:6633"
  is_connected: true
  Port "eth0"
    Interface "eth0"
  Port "sw1"
    Interface "sw1"
  Port "eth1"
    Interface "eth1"
  Port "ovsbr0"
    Interface "ovsbr0"
    type: internal
  Port "sw0"
    Interface "sw0"
  Port "eth2"
    Interface "eth2"
  Port "sw2"
    Interface "sw2"
  ovs_version: "2.9.8"
root@ovs1:/tmp/pycore.46249/ovs1.conf#

```

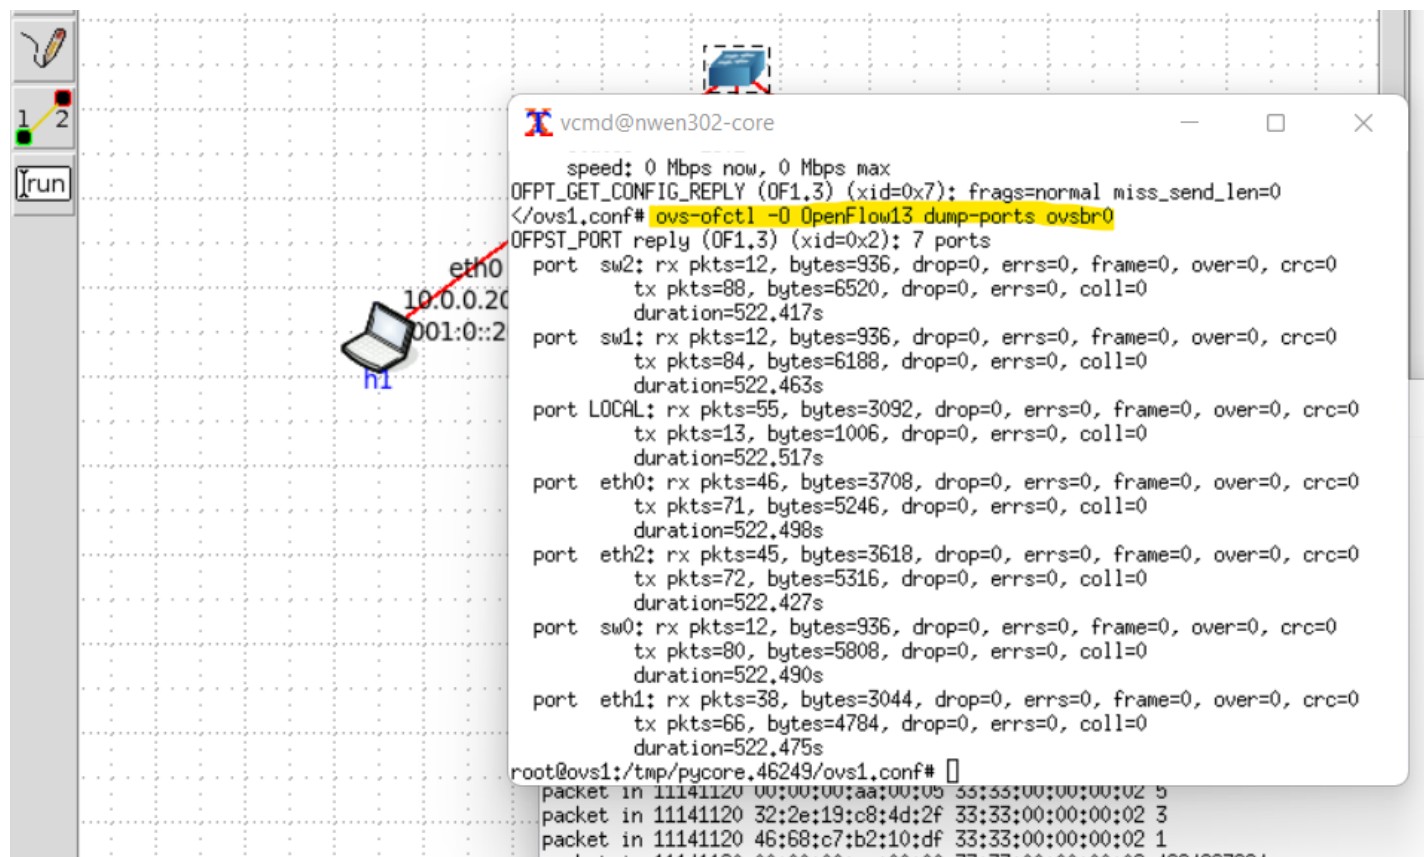


```

vcmd@nwen302-core
root@ovs1:/tmp/pycore.46249/ovs1.conf# ovs-ofctl -O OpenFlow13 show ovsbr0
OFPST_FEATURES_REPLY (OF1.3) (xid=0x2): dpid:0000000000aa0000
n_tables:254, n_buffers:0
capabilities: FLOW_STATS TABLE_STATS PORT_STATS GROUP_STATS QUEUE_STATS
OFPST_PORT_DESC reply (OF1.3) (xid=0x3):
1(eth0): addr:00:00:00:aa:00:00
  config: 0
  state: LIVE
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
2(sw0): addr:46:90:cc:fd:b8:59
  config: 0
  state: LIVE
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
3(eth1): addr:00:00:00:aa:00:02
  config: 0
  state: LIVE
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
4(sw1): addr:96:df:8a:22:ec:e2
  config: 0
  state: LIVE
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
5(eth2): addr:00:00:00:aa:00:04
  config: 0
  state: LIVE
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
6(sw2): addr:4e:23:7d:e3:40:69
  config: 0
  state: LIVE
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
LOCAL(ovsbr0): addr:00:00:00:aa:00:00
  config: 0
  state: LIVE
  speed: 0 Mbps now, 0 Mbps max
OFPST_GET_CONFIG_REPLY (OF1.3) (xid=0x7): frags=normal miss_send_len=0
root@ovs1:/tmp/pycore.46249/ovs1.conf#

```

The command 'ovs-ofctl -O OpenFlow13 dump-ports ovsbr0' prints the br0 OpenFlow ports statistics and shows detailed information about all the interfaces connected to this bridge, including the speed, state and peer information.



Now that I feel I understand Open vSwitches design and uses I will now be implementing the above pseudo code from the design section into program python code for the switch to run with to complete the given tasks.

KEY TASKS

Task 1

Modify `simple_switch_13.py` to include logic to **block** traffic between host h2 and host h3. Save the modified file as `lab3_task1.py`. Explain code and test procedures.

- To implement a blocking function I utilized the functions already provided in the SimpleSwitch13 class to include an if statement (after the code has learnt the src and dst values) that checks if the source or destination mac address is either from/to host 2 or host 3. Within the if statement I have included a print to terminal letting the user know the communication is blocked. If this statement is true after printing it will drop (by implementing an empty return). Lines 96 - 100.
- I got the source/destination MAC address information by inputting 'ifconfig' on each terminal, host 2 and host 3.
- Below is a screenshot of the code running through the switch and the hosts attempting to communicate with each other. (Shows that host 2 is able to ping host 1 but not host 3)

The screenshot displays a network topology and two terminal windows. The topology shows a central switch (labeled 'ovs1') connected to three hosts: h1, h2, and h3. Host h1 has IP 10.0.0.20/24 and MAC 00:1:0:20:64. Host h2 has IP 10.0.0.21/24 and MAC 20:01:0:20:64. Host h3 has IP 10.0.0.22/24 and MAC 20:01:2:20:64. The switch has interfaces eth0, eth1, and eth2 connected to h1, h2, and h3 respectively.

The left terminal window shows the output of the switch's log, indicating that traffic between h1 and h2 is successful, while traffic between h2 and h3 is blocked. The right terminal window shows the output of ping commands from h2 to h1 and h3, confirming that h2 can reach h1 but not h3.

```
es=254)
move onto main mode
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 06:b4:d8:b7:4f:fd 33:33:00:00:00:02 4
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 76:39:c8:92:8c:60 33:33:00:00:00:02 5
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:03 ff:ff:ff:ff:ff:ff 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:01 00:00:00:aa:00:03 1
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:03 00:00:00:aa:00:01 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:05 33:33:00:00:00:02 5
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:03 ff:ff:ff:ff:ff:ff 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:05 00:00:00:aa:00:03 5
blocked communication
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:03 33:33:00:00:00:02 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 36:11:e6:0d:00:ea 33:33:00:00:00:02 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:03 ff:ff:ff:ff:ff:ff 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:05 00:00:00:aa:00:03 5
blocked communication
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:00 33:33:00:00:00:02 4294967294
packet in 11141120 76:39:c8:92:8c:60 33:33:00:00:00:02 5
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 11141120 00:00:00:aa:00:03 ff:ff:ff:ff:ff:ff 3
```

```
root@h2:/tmp/pycore.45411/h2.conf# ping -c 2 10.0.0.20
PING 10.0.0.20 (10.0.0.20) 56(84) bytes of data:
64 bytes from 10.0.0.20: icmp_seq=1 ttl=64 time=10.1 ms
64 bytes from 10.0.0.20: icmp_seq=2 ttl=64 time=0.082 ms

--- 10.0.0.20 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1005ms
rtt min/avg/max/mdev = 0.082/5.116/10.150/5.034 ms
root@h2:/tmp/pycore.45411/h2.conf# ping -c 2 10.0.0.22
PING 10.0.0.22 (10.0.0.22) 56(84) bytes of data:
From 10.0.0.21 icmp_seq=1 Destination Host Unreachable
From 10.0.0.21 icmp_seq=2 Destination Host Unreachable

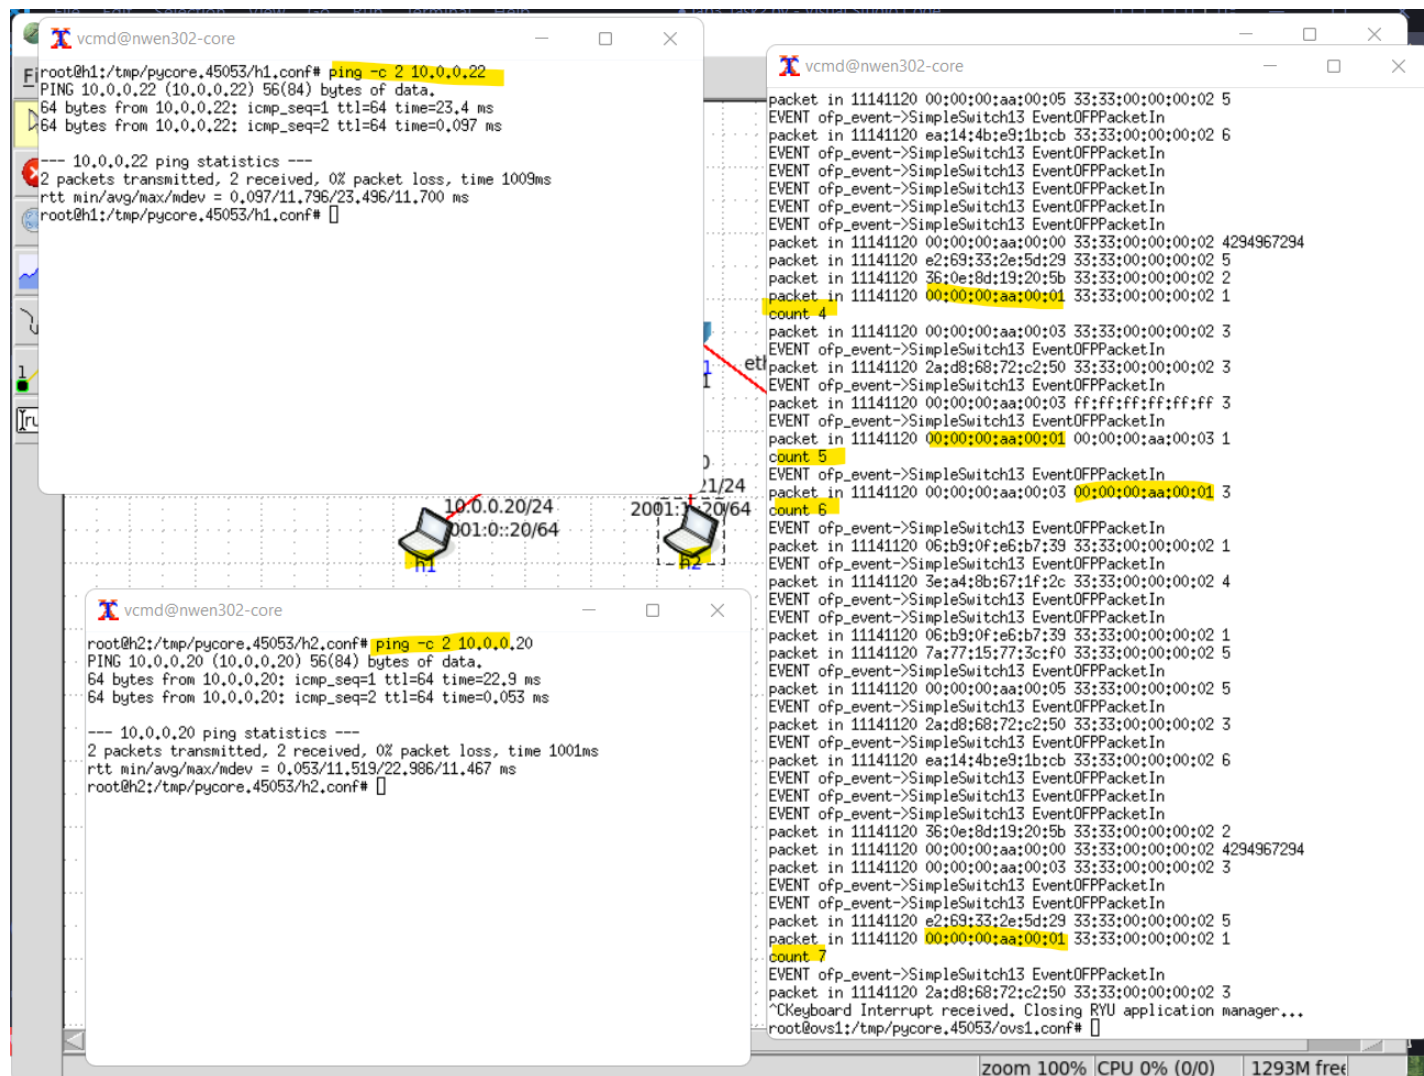
--- 10.0.0.22 ping statistics ---
2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 1022ms
pipe 2
root@h2:/tmp/pycore.45411/h2.conf#
```


Task 2

Modify simple_switch13.py to count all traffic going to and originating from host h1. Save the modified file as lab3_task2.py. Explain Code and test procedures.

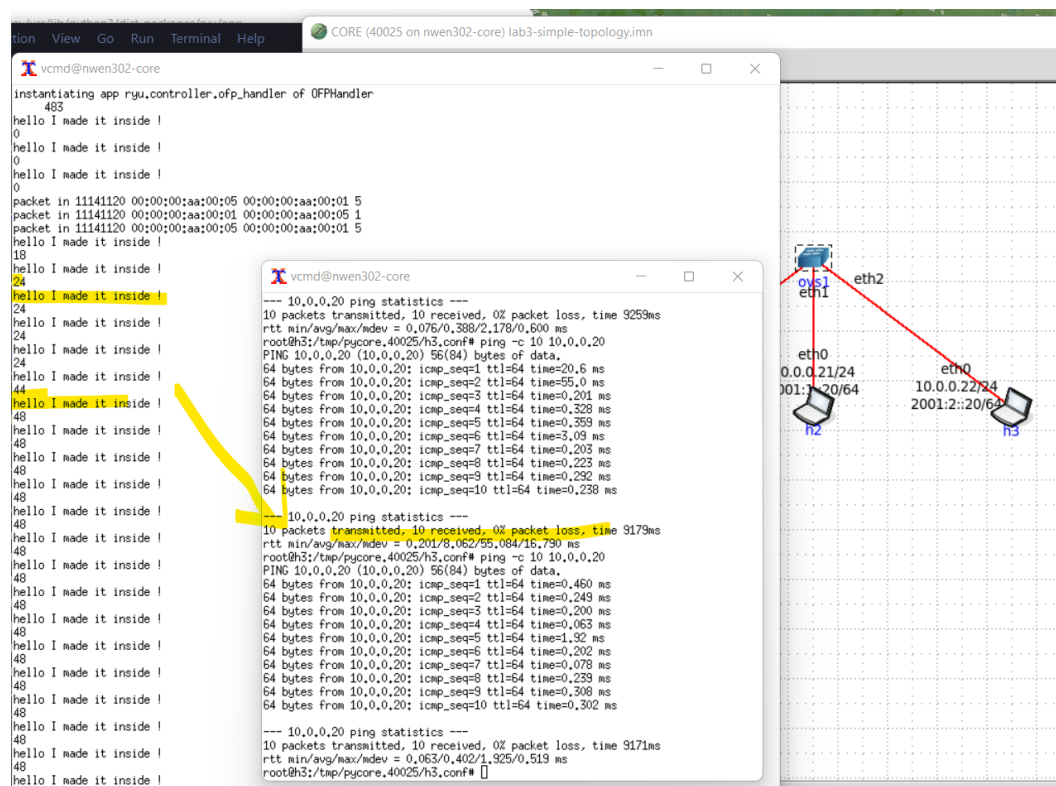
First implementation:

- To be able to count the traffic from a specified host I created a global num variable to use for keeping track of the traffic and updating it's value.
- To be able to access the num variable I just called it before my following if statement within the class itself.
- Similar to task 1, I implemented the src and dst variables and checked if either of them were from host 1 and if so, update the num variable (num + 1) and print the num value.
- Below is a screenshot of the executed code through the VM: (host 1 pinging host 3, host 2 pinging host 1. The switch terminal with the code execution shows that the traffic from host 1 is sufficiently being counted).



Second implementation:

- After some debugging I found that the count wasn't correct and had to implement my code another way. Referencing from the traffic_minotor library (shown within the references section). I had to implement a '_state_change_handler()', '_monitor()', '_request_stats()', and a '_port_stats_reply_handler'.
- Within the '_port_stats_reply_handler();' function I implemented a stats variable which is a dictionary that holds each host's traffic information. By accessing the first element of the stats dictionary (host 1) by stats[0].rx_packets + stats[0].tx_packets it returns the number value of all traffic that has been through host 1. Since this value holds the history of every single packet sent through host 1 the number is quite large as it doesn't count from our point of beginning the code.
- To debug this I created new variables where one will be counting the total packets and one that will be counting the 'new' packets minus the total packets. This will return only the packets that we initialize. From the screenshot below it shows my code running. The value begins at 0 before we ping host 1 which is good, then when we ping 10 packets to host 1 from host 2 the count value goes to 24. We account for 20 packets and the other 4 packets are the background broadcast that the program runs regardless of our code. When I ping 10 packets again the value goes from 24 to 48 which shows we successfully recorded the 20 send/receive packets communicating with host 1, then similar to the last time the Ryu is sending the background broadcast which adds another 4 packets to the count.

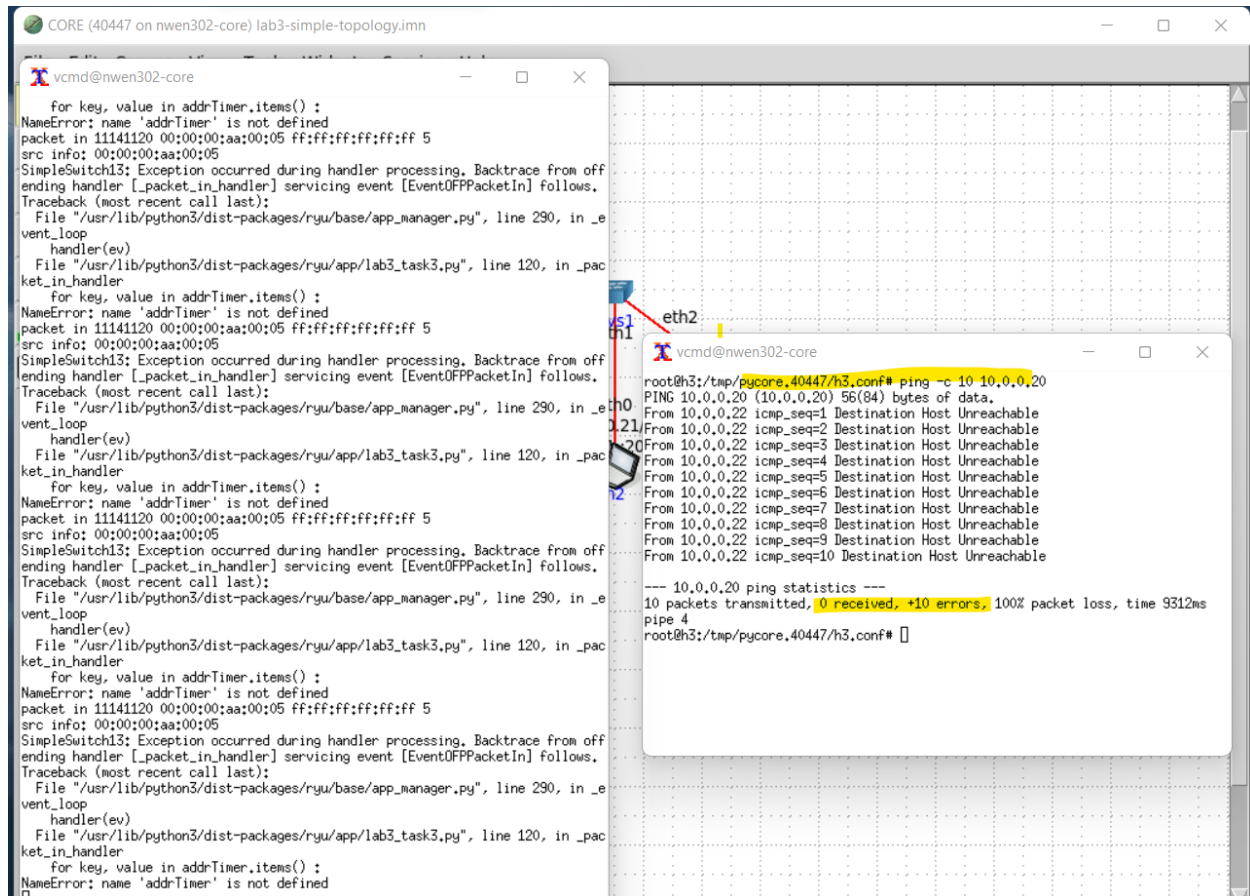


Task 3

Extend `simple_switch_13.py` to combine Task 1 and Task 2 functionalities. Keep track of all traffic (count the number of packets) originating from each host. If the counter exceeds a specific number, block all the traffic originating from this host for 1 minute. The maximum packet count number should be configured through the `MAX_COUNT` variable. Save the modified file as `lab3_task3.py`. Explain Code and test procedures.

First Implementation:

- Outside the class I created two separate 'dictionaries' (similar functioning to maps/arrays in java) that contain the hosts mac addresses. One map/dictionaries value holds an int value for counting and the other dictionaries value holds a boolean value which will be used for the blocking functionality later on.
- Imported the python threading library for opening the threads for each blocked host (each time a host is blocked it needs to be run on a new separate thread to not affect the program root processes).
- Thread function declaration at the beginning of the SimpleSwitch13 class. A simple method named `timerChecks` which passes in a `mac_address` and resets the `mac_address` from the `addrTimer` dictionary to false (false = unblocked, true = blocked). This gets called when the timer has hit 60 seconds as a means to reset the blocking.
- Added the hosts mac addresses to the `__init__` function as a map/dictionary to compare the src info to later on.
- Within the `_packet_in_handler` method the first thing I did was call the global variables/dictionaries I created and then did a simple if statement which checks the if the source (src) is inside the `mac_address` list (referring to the dictionary/map within the `__init__` function) and when it recognises it add the values count to '+1' on the specific source addresses element.
- For the timer threading function I created a for loop which goes through the keys and values of each mac address within the `addrTimer` dictionary and checks if the current element's source value is equal to true. If the value returns true, return, as a means to continue blocking the traffic.
- The last if statement is the counter function which checks if the source addresses value is equal to the `MAX_COUNT` variable (10) and if so, change the source addresses `addrTimer` value to true and open/start the thread which will count up to 60 seconds before running the `timerChecks` function which will unblock it.
- The way I have gone about testing this is after running the code through the switch terminal I have pinged from each host 10 packets to check if 1, the count sufficiently goes upwards and 2, when it hits 10 the communications are blocked and the terminal notifies us so.



- Above screenshots shows the testing output using the above logic

Second Implementation:

- Similar to the second implementation of task 2 I started with referencing from the traffic_minotar library and implemented a '_state_change_handler()', '_monitor()', '_request_stats()', and a '_port_stats_reply_handler'.
- Within the '_port_stats_reply_handler();' function I implemented a stats variable which is a dictionary that holds each host's traffic information.
- I did the same thing with task 2 where I checked if the stats[0] of rx_packets and tx_packets and then added the values to a count variable except I did so for each host with separate variables.
- Then I changed the hosts elements value within the addrCount variable to increment by 1 as a means to count upwards.
- Then I made a loop to go through the keys and values of the addrCount to assess. When the hosts value within the addrCount dictionary reaches the MAX_COUNT (10) it will then start a timing thread which will block traffic through that host for 60 seconds. After the 60seconds the timerchecks()function will be called which will unblock the traffic.

Note:

- Through testing I believe I was close but unfortunately did not get the results I had hoped for. The screenshots below show the output I was getting, It appears it was somewhat counting (the number 0 to 26 is from host 3 pinging host 1 10 times, so 20 packets went to and from host 1) but unsure if the blocking is successful as I attempted to ping 10 packets more than once. What should have happened was after host 1's count value hit 10 the manual pings from host 3 would no longer be successful.

```
vcmd@nwen302-core lab3-simple-topology.imn
actions = [parser.OFPActionOutput(out_port)]
NameError: name 'addrTimer' is not defined
packet in 11141120 fe:9e:fa:a8:1c:14 33:33:00:00:00:02 1
src info: fe:9e:fa:a8:1c:14
SimpleSwitch13: Exception occurred during handler processing. Backtrace from off
ending handler [_packet_in_handler] servicing event [EventOFPacketIn] follows.
Traceback (most recent call last):
  File "/usr/lib/python3/dist-packages/ryu/base/app_manager.py", line 290, in _e
vent_loop
    handler(ev)
  File "/usr/lib/python3/dist-packages/ryu/app/lab3_task3.py", line 120, in _pac
ket_in_handler
    actions = [parser.OFPActionOutput(out_port)]
NameError: name 'addrTimer' is not defined
packet in 11141120 00:00:00:aa:00:00 33:33:00:00:00:02 4294967294
src info: 00:00:00:aa:00:00
SimpleSwitch13: Exception occurred during handler processing. Backtrace from off
ending handler [_packet_in_handler] servicing event [EventOFPacketIn] follows.
Traceback (most recent call last):
  File "/usr/lib/python3/dist-packages/ryu/base/app_manager.py", line 290, in _e
vent_loop
    handler(ev)
  File "/usr/lib/python3/dist-packages/ryu/app/lab3_task3.py", line 120, in _pac
ket_in_handler
    actions = [parser.OFPActionOutput(out_port)]
NameError: name 'addrTimer' is not defined
<usr/lib/python3/dist-packages/ryu/app/lab3_task3.py
Registered VCS backend: git
Registered VCS backend: hg
Registered VCS backend: svn
Registered VCS backend: bzr
loading app /usr/lib/python3/dist-packages/ryu/app/lab3_task3.py
loading app ryu.controller.ofp_handler
instantiating app /usr/lib/python3/dist-packages/ryu/app/lab3_task3.py of Simple
Switch13
instantiating app ryu.controller.ofp_handler of OFPHandler
SimpleSwitch13: Exception occurred during handler processing. Backtrace from off
ending handler [_state_change_handler] servicing event [EventOFPPStateChange] fol
lows.
Traceback (most recent call last):
  File "/usr/lib/python3/dist-packages/ryu/base/app_manager.py", line 290, in _e
vent_loop
    handler(ev)
  File "/usr/lib/python3/dist-packages/ryu/app/lab3_task3.py", line 144, in _sta
te_change_handler
    if datapath.id not in self.datapaths:
AttributeError: 'SimpleSwitch13' object has no attribute 'datapaths'
packet in 11141120 00:00:00:aa:00:05 ff:ff:ff:ff:ff:ff 5
packet in 11141120 00:00:00:aa:00:01 00:00:00:aa:00:05 1
packet in 11141120 00:00:00:aa:00:05 00:00:00:aa:00:01 5
```

```
vcmd@nwen302-core
From 10.0.0.22 icmp_seq=8 Destination Host Unreachable
From 10.0.0.22 icmp_seq=9 Destination Host Unreachable
From 10.0.0.22 icmp_seq=10 Destination Host Unreachable
--- 10.0.0.20 ping statistics ---
10 packets transmitted, 0 received, +10 errors, 100% packet loss, time 9312ms
pipe 4
root@h3:/tmp/pycore.40447/h3.conf# ping -c 10 10.0.0.20
PING 10.0.0.20 (10.0.0.20) 56(84) bytes of data:
64 bytes from 10.0.0.20: icmp_seq=1 ttl=64 time=0.265 ms
64 bytes from 10.0.0.20: icmp_seq=2 ttl=64 time=0.200 ms
64 bytes from 10.0.0.20: icmp_seq=3 ttl=64 time=0.199 ms
64 bytes from 10.0.0.20: icmp_seq=4 ttl=64 time=0.096 ms
64 bytes from 10.0.0.20: icmp_seq=5 ttl=64 time=0.241 ms
64 bytes from 10.0.0.20: icmp_seq=6 ttl=64 time=0.215 ms
64 bytes from 10.0.0.20: icmp_seq=7 ttl=64 time=0.198 ms
64 bytes from 10.0.0.20: icmp_seq=8 ttl=64 time=0.326 ms
64 bytes from 10.0.0.20: icmp_seq=9 ttl=64 time=0.199 ms
64 bytes from 10.0.0.20: icmp_seq=10 ttl=64 time=0.164 ms
--- 10.0.0.20 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9194ms
rtt min/avg/max/mdev = 0.096/2.837/26.534/7.899 ms
root@h3:/tmp/pycore.40447/h3.conf#
```

- The above screenshots show the terminal output from running my updated task3 code using the above implementation. I believe I was almost there but did not quite execute it.

References

Python Tutorial,

- <https://docs.python.org/3/tutorial/>

Ryu Tutorial & Documentation,

- https://ryu.readthedocs.io/en/latest/writing_ryu_app.html

Ryu Packet Library,

- https://ryu.readthedocs.io/en/latest/library_packet.html

Traffic Monitor Code,

- https://osrg.github.io/ryu-book/en/html/traffic_monitor.html

Appendices

1. Appendix A - Software

Windows Linux Subsystem, WSL & Ubuntu

<https://ubuntu.com/tutorials/install-ubuntu-on-wsl2-on-windows-11-with-gui-support#1-overview>

WinSCP 5.21.2, Windows 64bit

<https://winscp.net/eng/download.php>

PuTTY 0.77, Windows 64bit

<https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html>

VirtualBox VM 6.1.38, Windows 64bit

<https://www.virtualbox.org/wiki/Downloads>

2. Appendix B - Configurations

Ryu SDN Framework

<https://ryu-sdn.org/>

Open vSwitch

<https://www.openvswitch.org/>

3. Appendix B - Configurations

Ryu SDN Framework

<https://ryu-sdn.org/>

Open vSwitch

<https://www.openvswitch.org/>

Acknowledgements

- Tutor, Zach Kingsford, with coding and development help
- Professors, Winston Seah and Alvin Valera, for development and general networking knowledge